

HDFS



Basics – *optional*

- Program
 - sequence of instructions written to perform a specified task with a computer
 - or a piece of code
- Process
 - an instance of a computer program that is being executed.
 - or a execution of a program
- Daemon Process
 - process which runs in background and has no controlling terminal.
- JVM – Java Virtual Machine
 - program which executes certain programs, namely those containing Java bytecode instructions

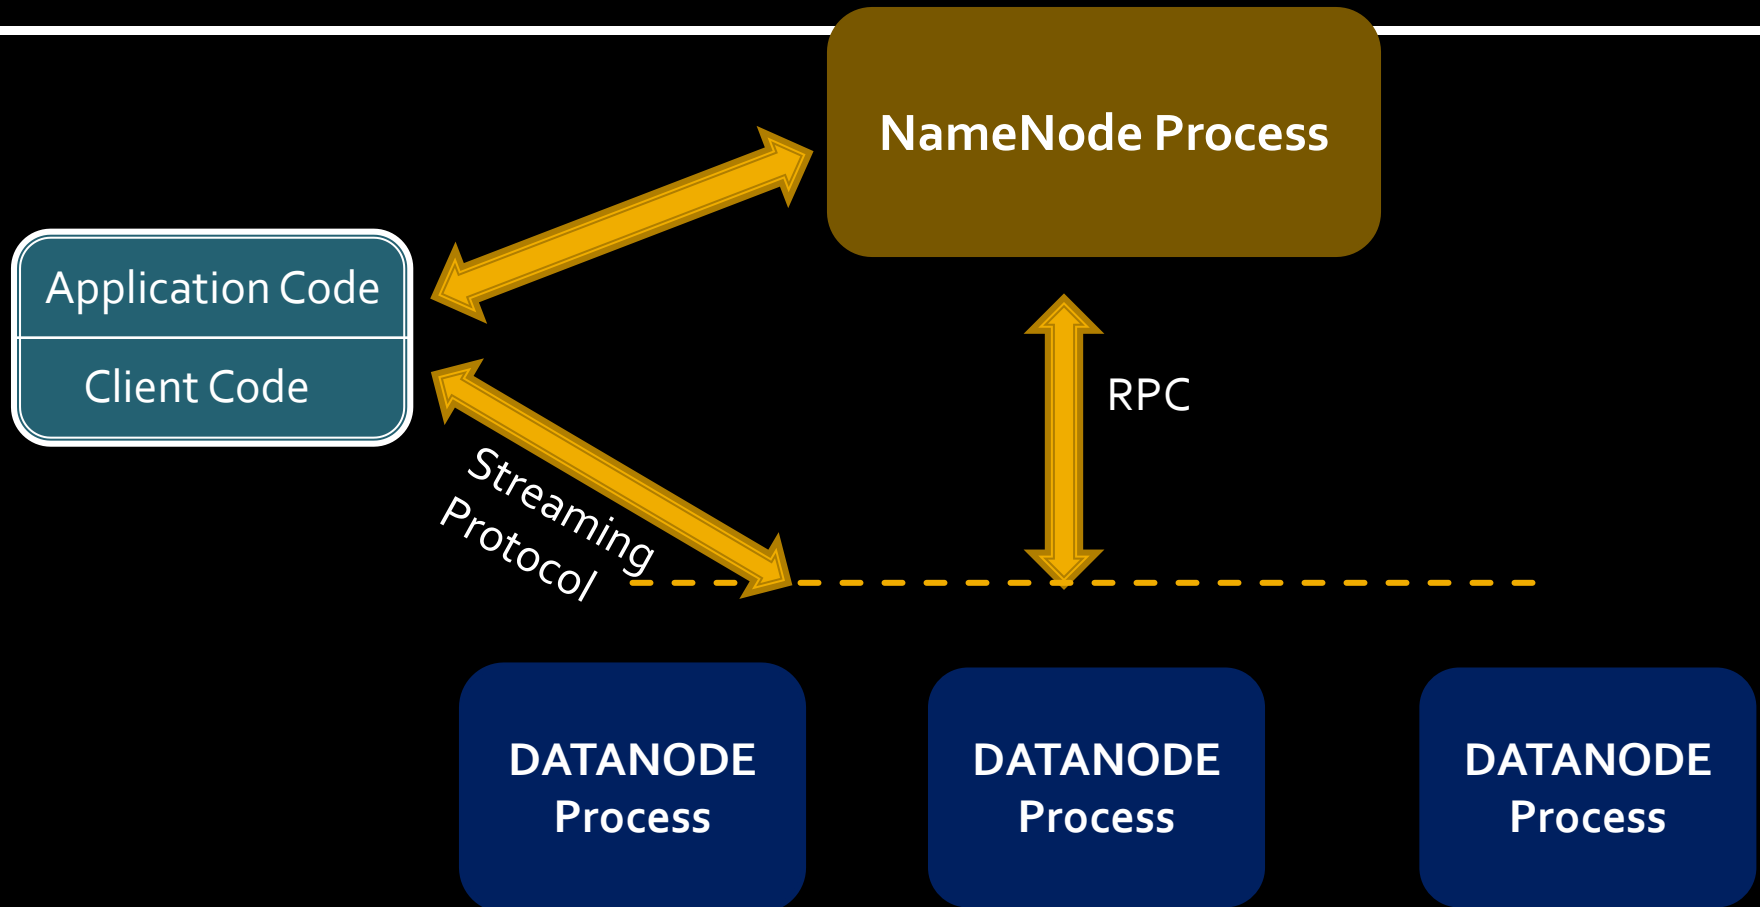
Basics – optional

- Client-server Concept
 - Client sends requests to one or more servers which in turn accepts, processes them and return the requested information to the client.
 - A server might run a software which listens on particular ip and port number for requests
 - Examples:
 - Server - web server
 - Client – web browser

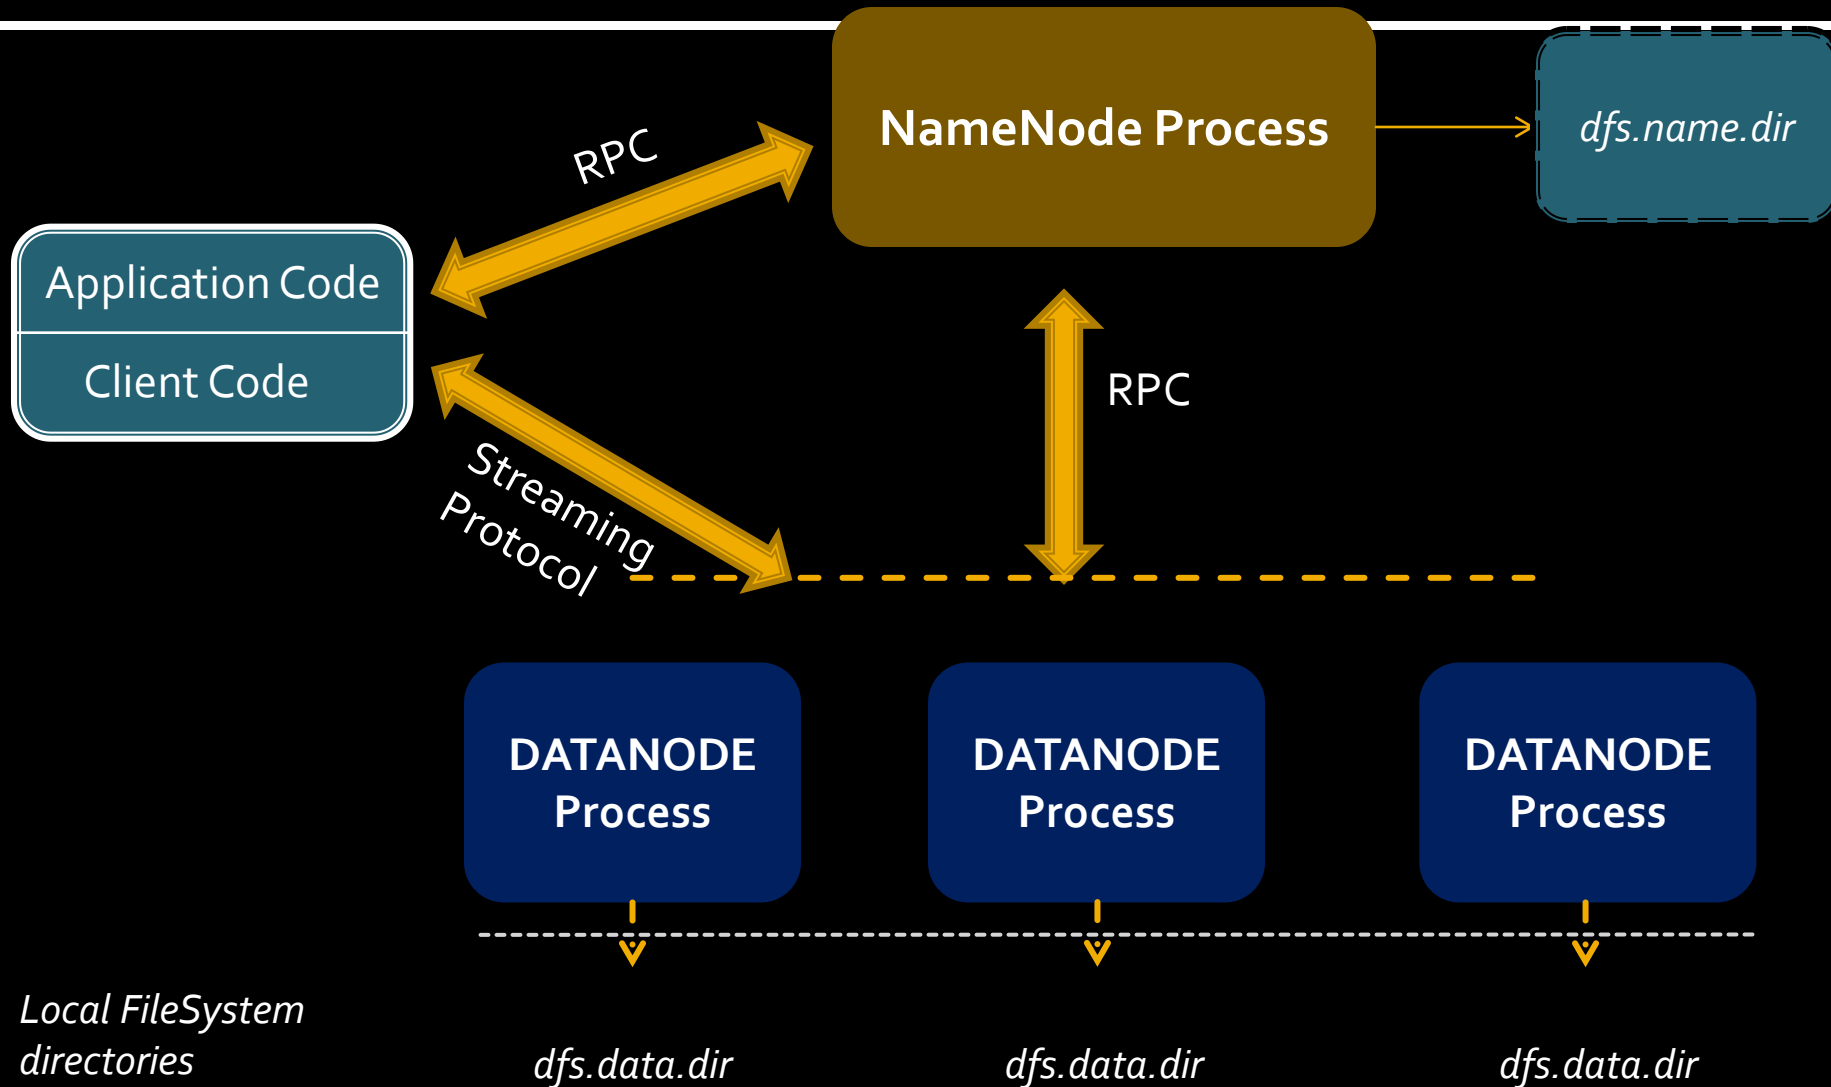
Introduction

- A distributed File System - **STORAGE**
 - A File System on multiple machines which sits on native filesystem
 - ext4,ext3
 - Hardware Failure
 - Due to usage of Commodity machines, failure is a common phenomenon
 - Designed for failure
 - Large Data Sets
 - Small Files Problem Due to NameNode
 - Simple Coherency Model
 - Write Once , Read Many Times
 - Streaming Data Access
 - High Throughput instead of low latency access
- ext3? ext4?

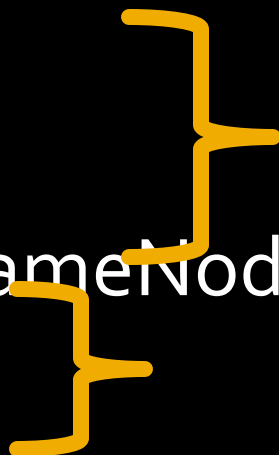
Continued...



Continued...



Daemons in Hadoop Core

- NameNode
 - DataNode
 - Secondary NameNode*
 - JobTracker*
 - TaskTracker*
- HDFS
- MR
- 

* - will be seen later in class

Block Concept

TestFile1.txt -> 1GB
Block Size -> 64 MB

Files are splitted into number of
chunks(Blocks) of pre-defined
size

No of Blocks = $1\text{GB} / 64\text{MB} = 16$ blocks
Blocks are B₁,B₂,.....B₁₆

DataNode

DataNode

DataNode

DataNode

B₁
B₈
B₁₀
B₁₃

B₃
B₇
B₁₂
B₁₆

B₄
B₅
B₁₁
B₁₄

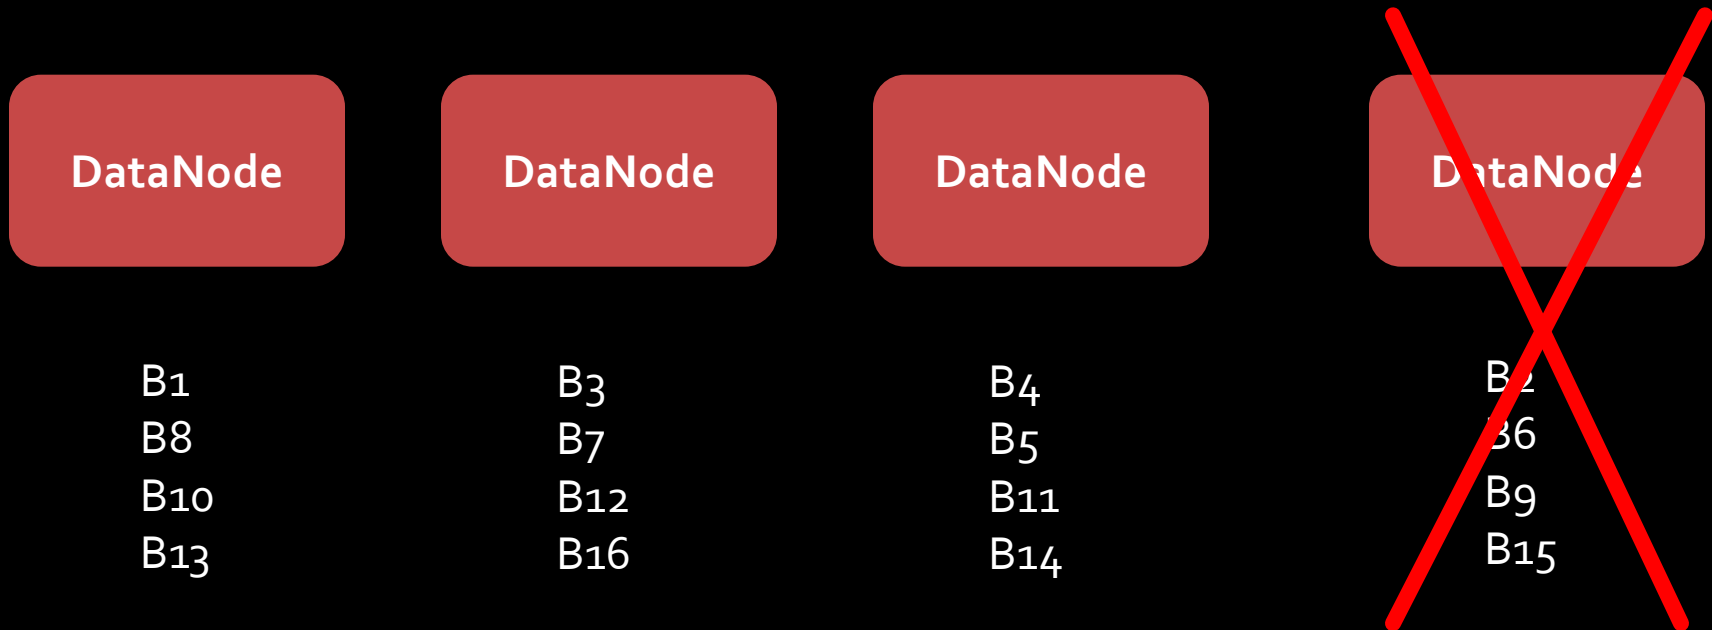
B₂
B₆
B₉
B₁₅

Block Concept

What happens to my data if
node 4 goes down??

TestFile1.txt -> 1GB
Block Size -> 64 MB

No of Blocks = $1\text{GB} / 64\text{MB} = 16$ blocks
Blocks are B₁, B₂, B₁₆

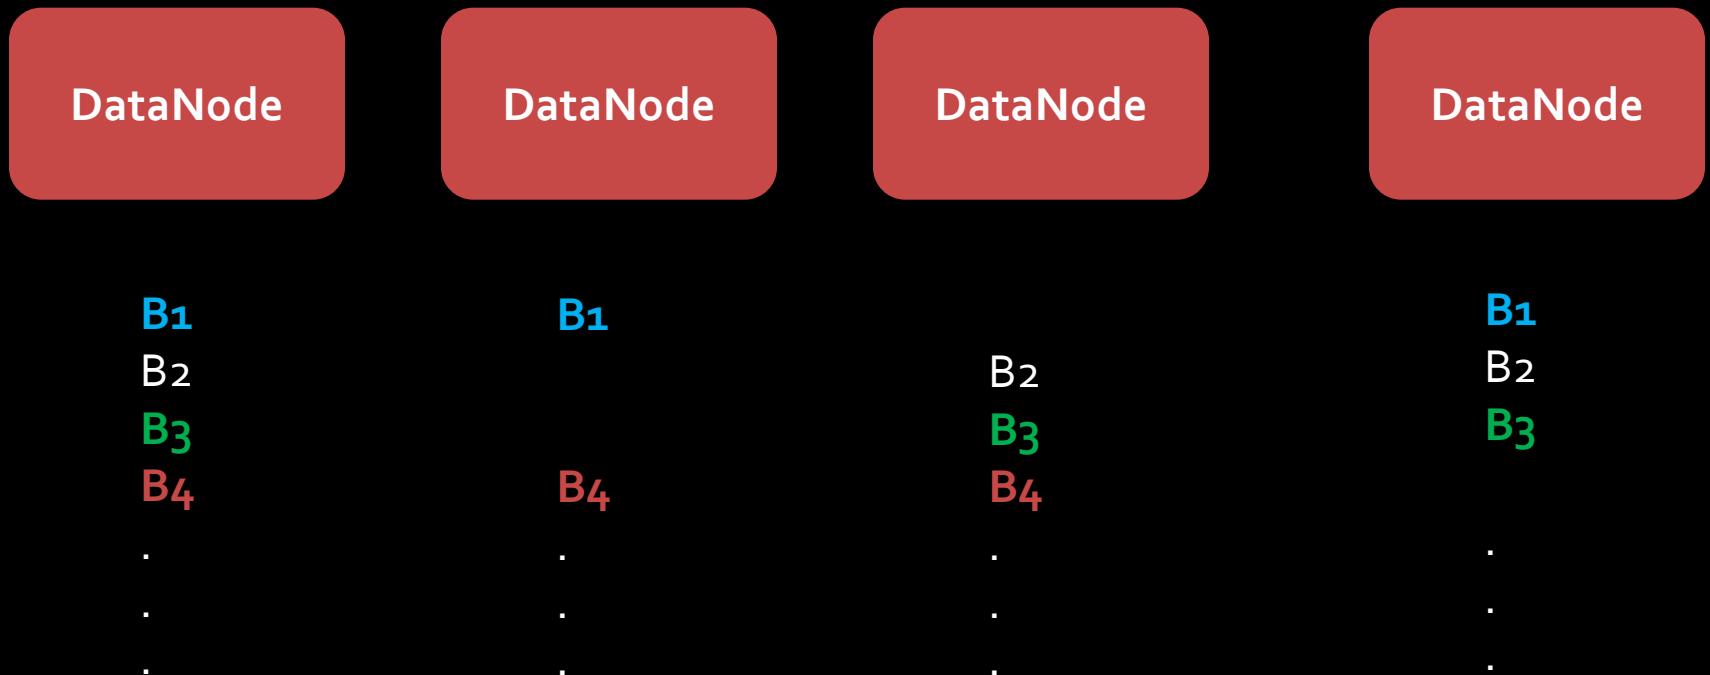


Fault Tolerant in HDFS

HDFS provides fault tolerant by
replication of each block by 3

TestFile1.txt -> 1GB
Block Size -> 64 MB

No of Blocks = $1\text{GB} / 64\text{MB} = 16$ blocks
Blocks are B₁, B₂, B₁₆



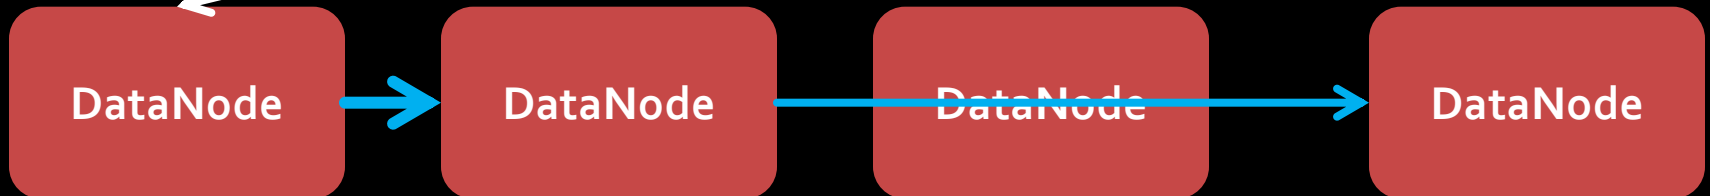
Data Pipelining

TestFile1.txt -> 1GB
Block Size -> 64 MB

No of Blocks = $1\text{GB} / 64\text{MB} = 16$ blocks
Blocks are B₁, B₂, B₁₆

Client

Write first block



B₁
B₂
B₃
B₄

.
. .
.

B₁

B₄

.
. .
.

B₂
B₃
B₄

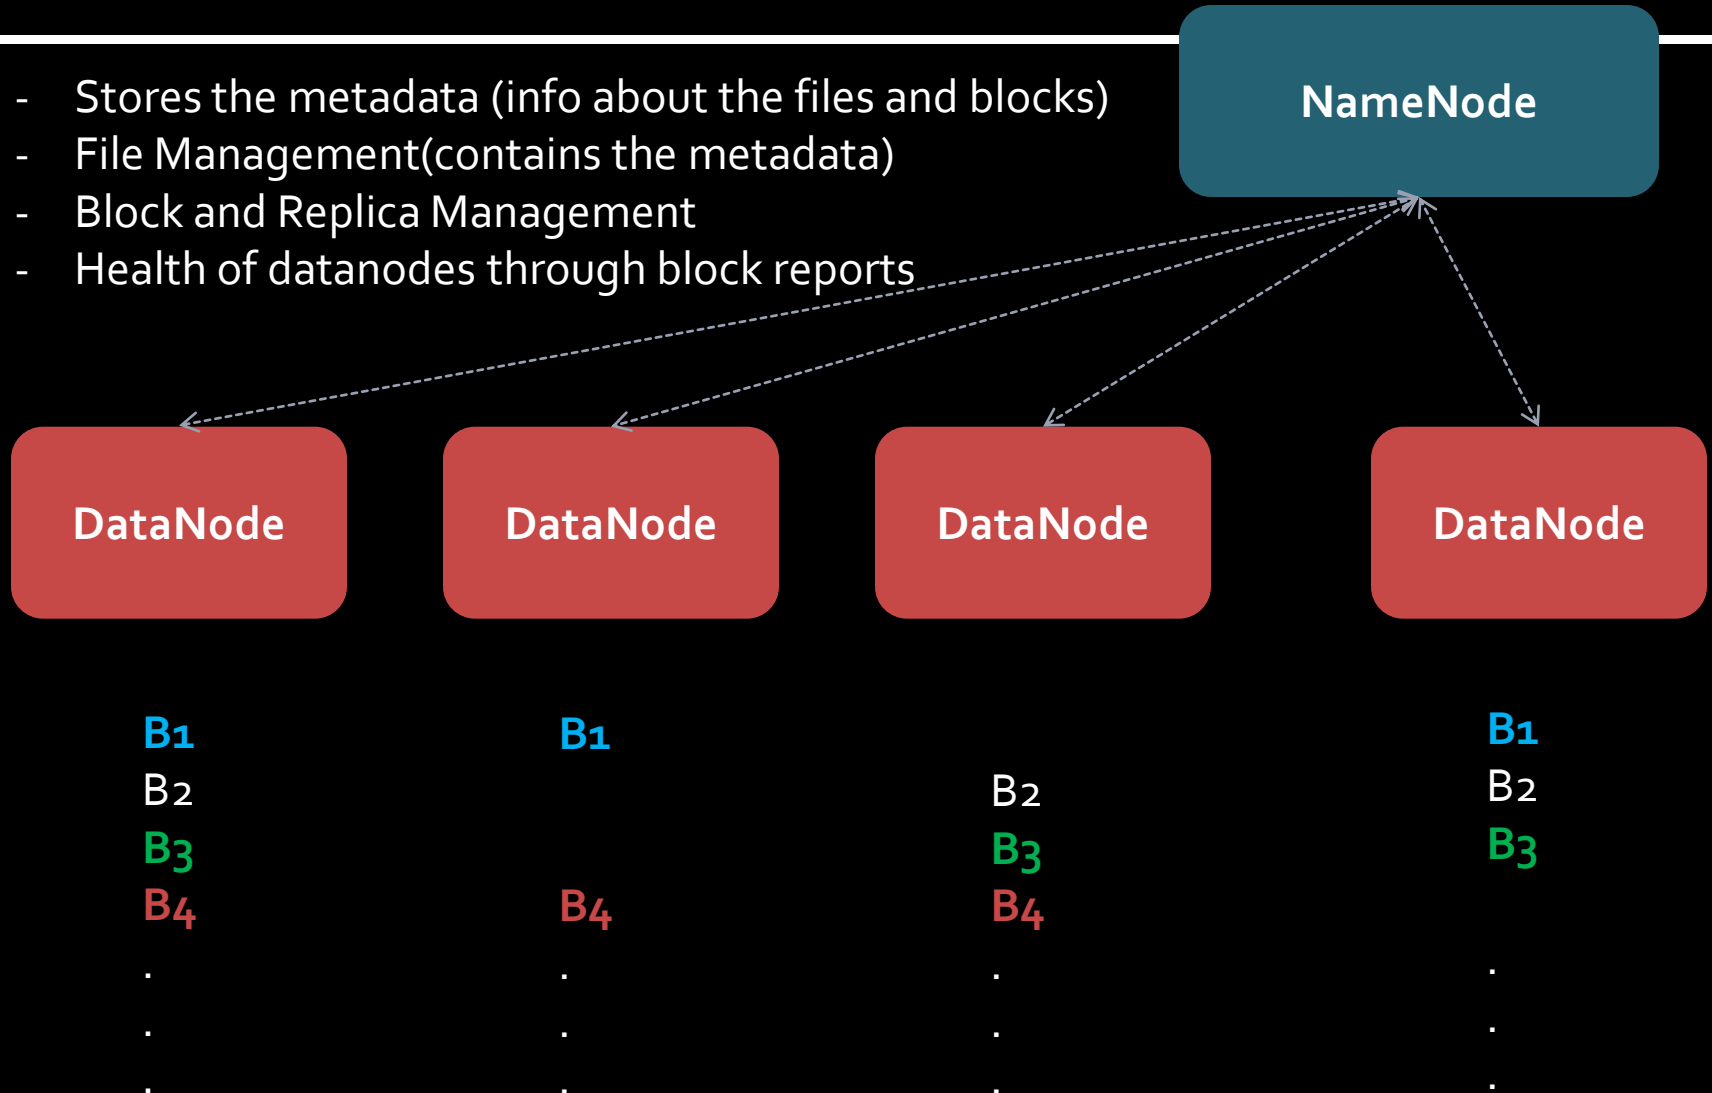
.
. .
.

B₁
B₂
B₃

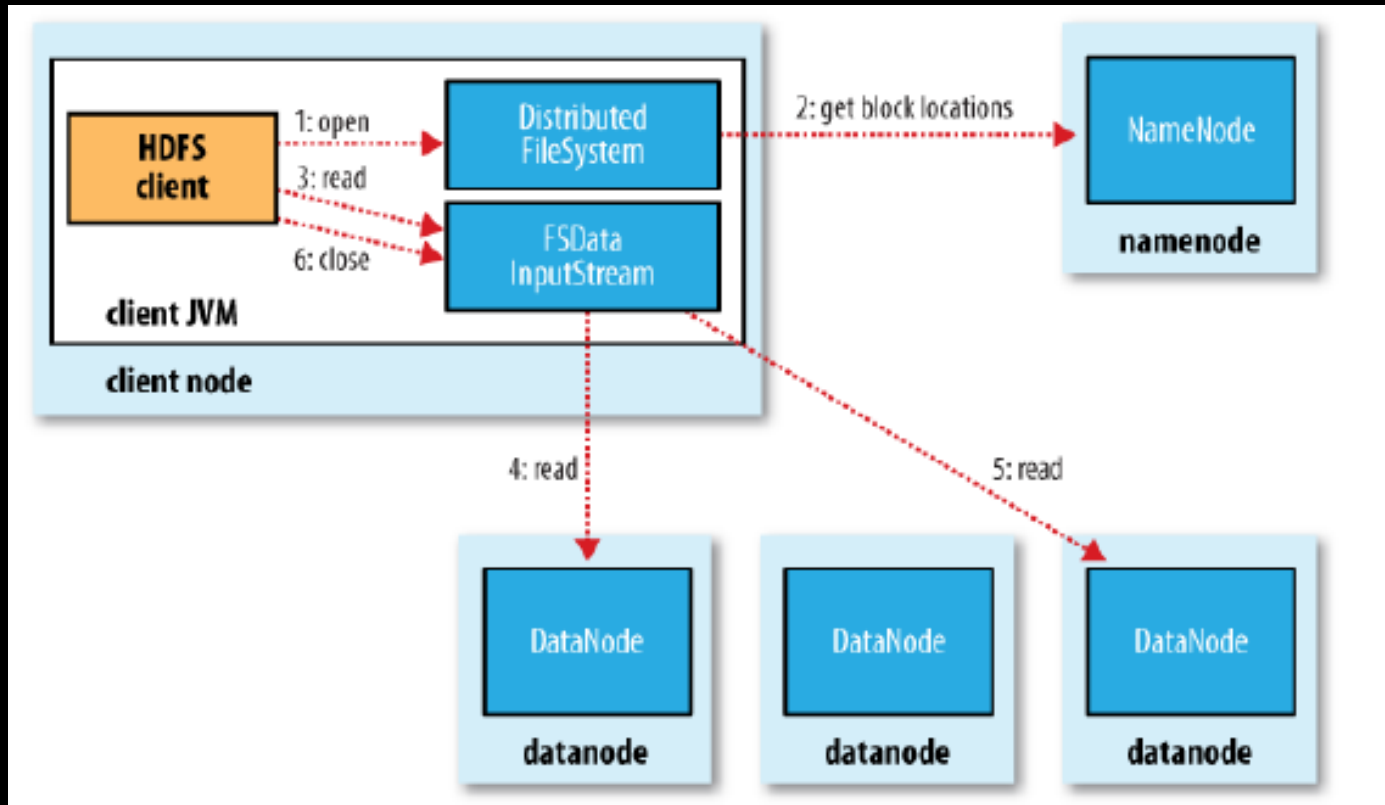
.
. .
.

Role of NameNode

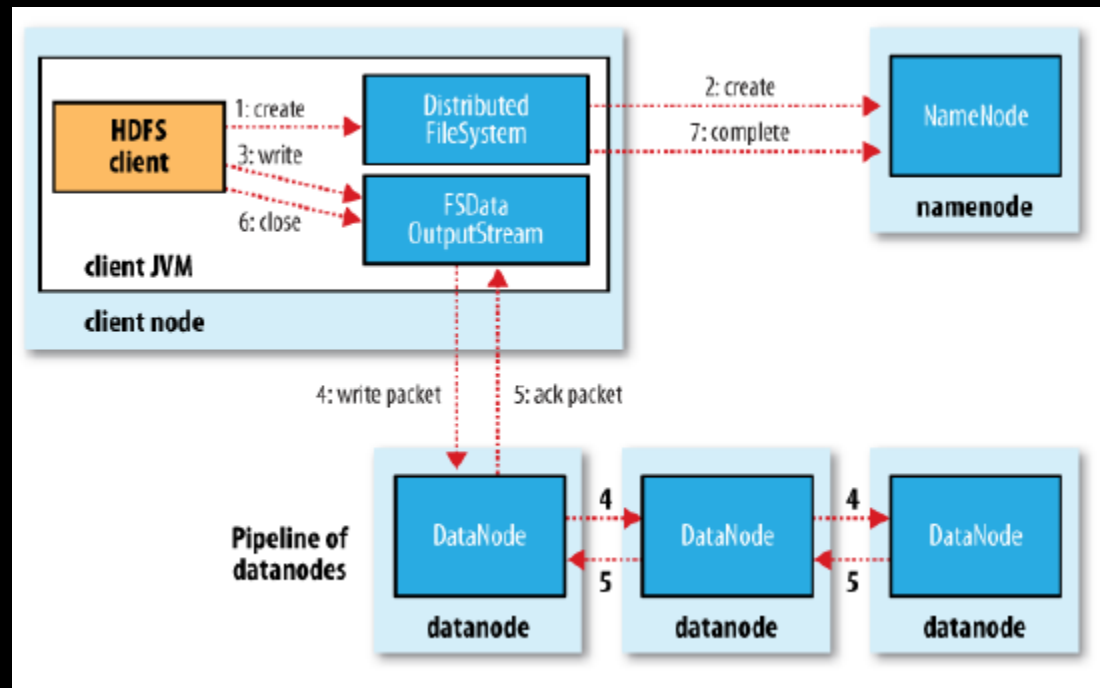
- Stores the metadata (info about the files and blocks)
- File Management(contains the metadata)
- Block and Replica Management
- Health of datanodes through block reports



File Read - Flow



File Write - Flow



Execution Modes & Installation

- Modes
 - Single Stand Alone
 - All Process runs in a single jvm
 - Does not use HDFS
 - Pseudo Distributed Mode - **for our training**
 - All daemon process runs in separate jvm in a single local machine
 - Used for development and testing
 - Uses HDFS to store data
 - Distributed Mode
 - A cluster of nodes more than 1
 - Each Process may run in different nodes
- Please follow instructor and doc provided

Installation

- Please follow the steps in the document given

After installation

- jps
 - Jps - jvm profiling status tool
- Web UI
 - NameNode - <http://localhost:50070>
 - JobTracker – <http://localhost:50030>

Accessing HDFS

- Command line
 - Usage: *hadoop dfs <command>*
- JAVA API
- webHDFS

HDFS commands

- *hadoop dfs -copyFromLocal <srcLOCALfile> <destHDFSfile>*
- *hadoop dfs -ls /*
- *hadoop dfs -cat /<destHDFSfile>*
- *hadoop dfs -copyToLocal <srcHDFSfile> <destLOCALfile>*
- *hadoop dfs -mkdir /test*
- *hadoop dfs -rmr /test*

JAVA API

- Most Packages Used
 - *org.apache.hadoop.conf.Configuration*
 - *org.apache.hadoop.fs.BlockLocation*
 - *org.apache.hadoop.fs.FSDataInputStream*
 - *org.apache.hadoop.fs.FSDataOutputStream*
 - *org.apache.hadoop.fs.FileStatus*
 - *org.apache.hadoop.fs.FileSystem*
 - *org.apache.hadoop.fs.Path*
 - *org.apache.hadoop.hdfs.DistributedFileSystem*
 - *org.apache.hadoop.hdfs.protocol.DataTransferInfo*

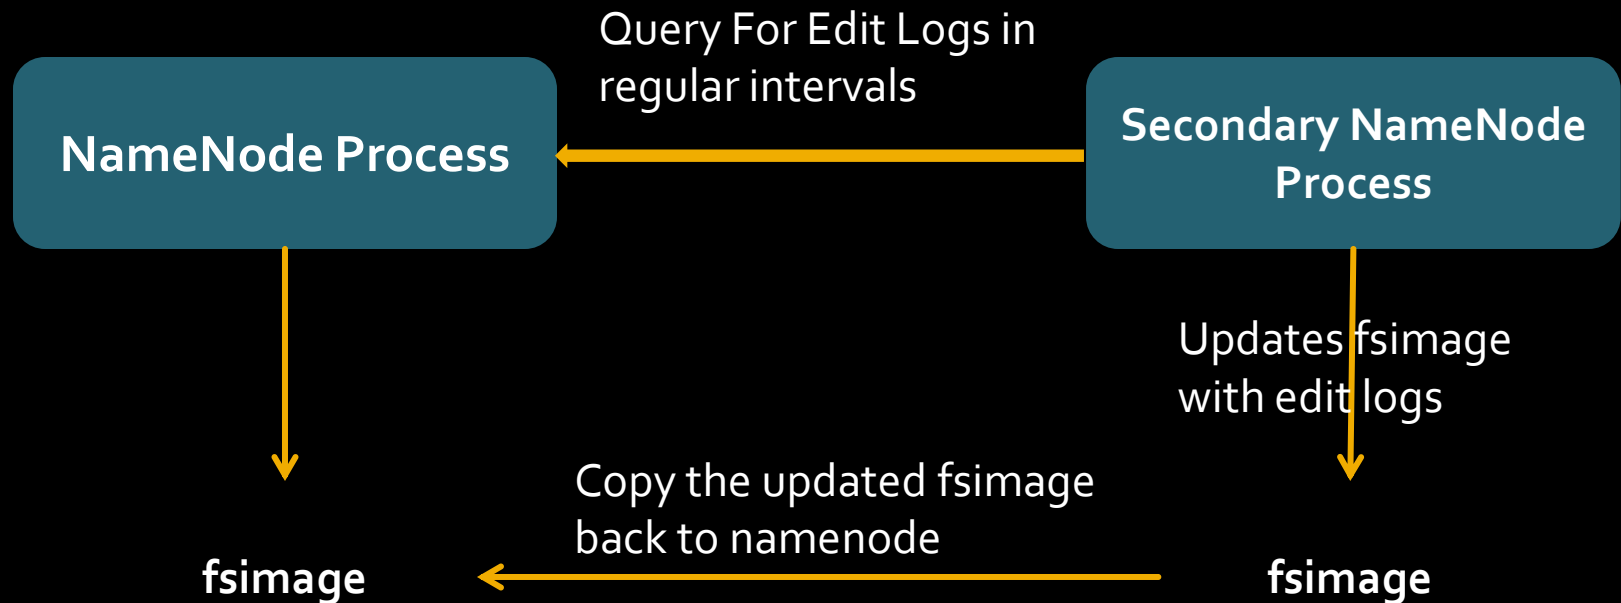
FileSystem API methods

- `append()`
- `copyFromLocalFile()`
- `create()`
- `delete()`
- `mkdirs()`
- `open()`

Secondary NameNode

- A helper node for NameNode
- Performs memory-intensive administrative functions for the NameNode
- Have a checkpoint for the file system (HDFS)
- Not a Backup Node

Role of Secondary NameNode



THANK YOU!!