

MAP REDUCE



Introduction

- A High level abstracted framework for distributed processing of large datasets
 - Fault Tolerant , Parallelization
- Computation consists of two phases
 - Map
 - Reduce
- A Master-Slaves architecture
- Computations occurs in multiple slave nodes
- And it tries to provide data locality as much as possible.

Daemons

■ JobTracker

- Client submits the computation to JobTracker
- Assign a task to the TaskTracker who has free slots and where data is stored if possible
 - It tries to provide data locality as much as possible.

■ TaskTracker

- Spawns a JVM process for each input split as directed by Job Tracker
- Send periodic heartbeats to Job Tracker

Terminology

- Job
 - A complete user defined computation or program
- Tasks
 - A subset of computation
 - Can be either execution of MAP or REDUCE
- Task Attempt
 - An attempt to run a task.
 - If an attempt fails, Job Tracker tries to start another task attempt for the same task.
 - By Default, total number of task attempts for a task is four

Anatomy of MR Code

- Mapper - a Java class to be extended by the developer
 - Methods – setup, map, run, cleanup
 - Map method takes a key value and can emit zero or more intermediate key value pairs depending upon the logic implemented by the developer
 - A JVM running Mapper is launched for each input split.
- Reducer – a Java class to be extended by the developer
 - Methods – setup, reduce, run, cleanup
 - Reduce method takes a (intermediate key-list of values) and can emit zero or more key value pairs depending upon the logic implemented by the developer
- Driver
 - Configures the job and submits the job to the cluster from the client.

Writables

■ *What is Serialization??*

- *the process of converting a data structure or object state into a format that can be stored (for example, in a file or memory buffer, or transmitted across a network connection link) and "resurrected" later in the same or another computer environment*

■ *Writables in Hadoop are present for serialization*

- *Why a separate framework instead of java serialization?*
 - *Compact*
 - *Fast*
 - *Extensible*
 - *Interoperable*

JAVA	Writable implementation	Serialized Size (bytes)
byte	ByteWritable	1
Boolean	BooleanWritable	1
Int	IntWritable	4
	VIntWritable	1-5
Float	FloatWritable	4
Double	DoubleWritable	8

NullWritable-----> usage NullWritable.get()

ArrayWritable

ArrayPrimitiveWritable

TwoDArrayWritable


MapWritable --- used widely..... Look @ Map or Hash

Map

SortedMapWritables

EnumSetWritable

Block Vs Split



Partitioner

- It decides which key (with its associated value) goes to reducer.
- By default its HashPartitioner.
 - hashes a record's key to determine which partition (and thus which reducer) the record belongs in

```
public class HashPartitioner<K, V> extends Partitioner<K, V> {  
  
    public int getPartition(K key, V value, int numReduceTasks) {  
        return (key.hashCode() & Integer.MAX_VALUE) % numReduceTasks;  
    }  
}
```

- Number of Partitions is equal to the number of reducers.

Custom Partitioner

- Need Custom partitioner for balancing(performance)
- To Write a partitioner, follow the steps
 - Extend **partitioner** class
 - Override the method **getPartition**
 - input – key,value,number of Reducers
 - Output –0 to n-1(where n - number of reducer)
 - `job.setPartitionerClass(<yourclassname.class>)`

COMBINER

- To Reduce the intermediate data from mapper to reducer
- To reduce Network IO and Disk IO
- Runs on a single mapper output(like a mini-reducer)
- Extends the Reducer class (new API)
- `Job.setCombinerClass(*.class);`
- Combiner may or maynot run
- Its better to use identical combiner and reduce when commutative as well as associative functions.

WritableComparable

- A WritableComparable is a Writable which is also Comparable
 - Two WritableComparables can be compared against each other to determine their 'order'
 - Keys must be WritableComparables because they are passed to the Reducer in sorted order
 - We will talk more about WritableComparable later

IntWritable

LongWritable

Text

WritableComparable

Writable

Defines a sort order. All keys must be WritableComparable

Defines a de/serialization protocol. Every data type in Hadoop is a Writable

WritableComparable

WritableComparable is a sub-interface of Writable

- Writable is an interface and must inherit following methods
 - ✓ `readFields(DataInput in);`
 - ✓ `write(DataOutput out);`
- Must implement `compareTo`, `hashCode`, `equals` methods
- All keys in MapReduce must be WritableComparable
- `compareTo` method compares the keys in the mapper out to provide sorted
 - ✓ It deserializes the keys and compare the values in java
 - ✓ It is time consuming .instead you can use comparators

Streaming API

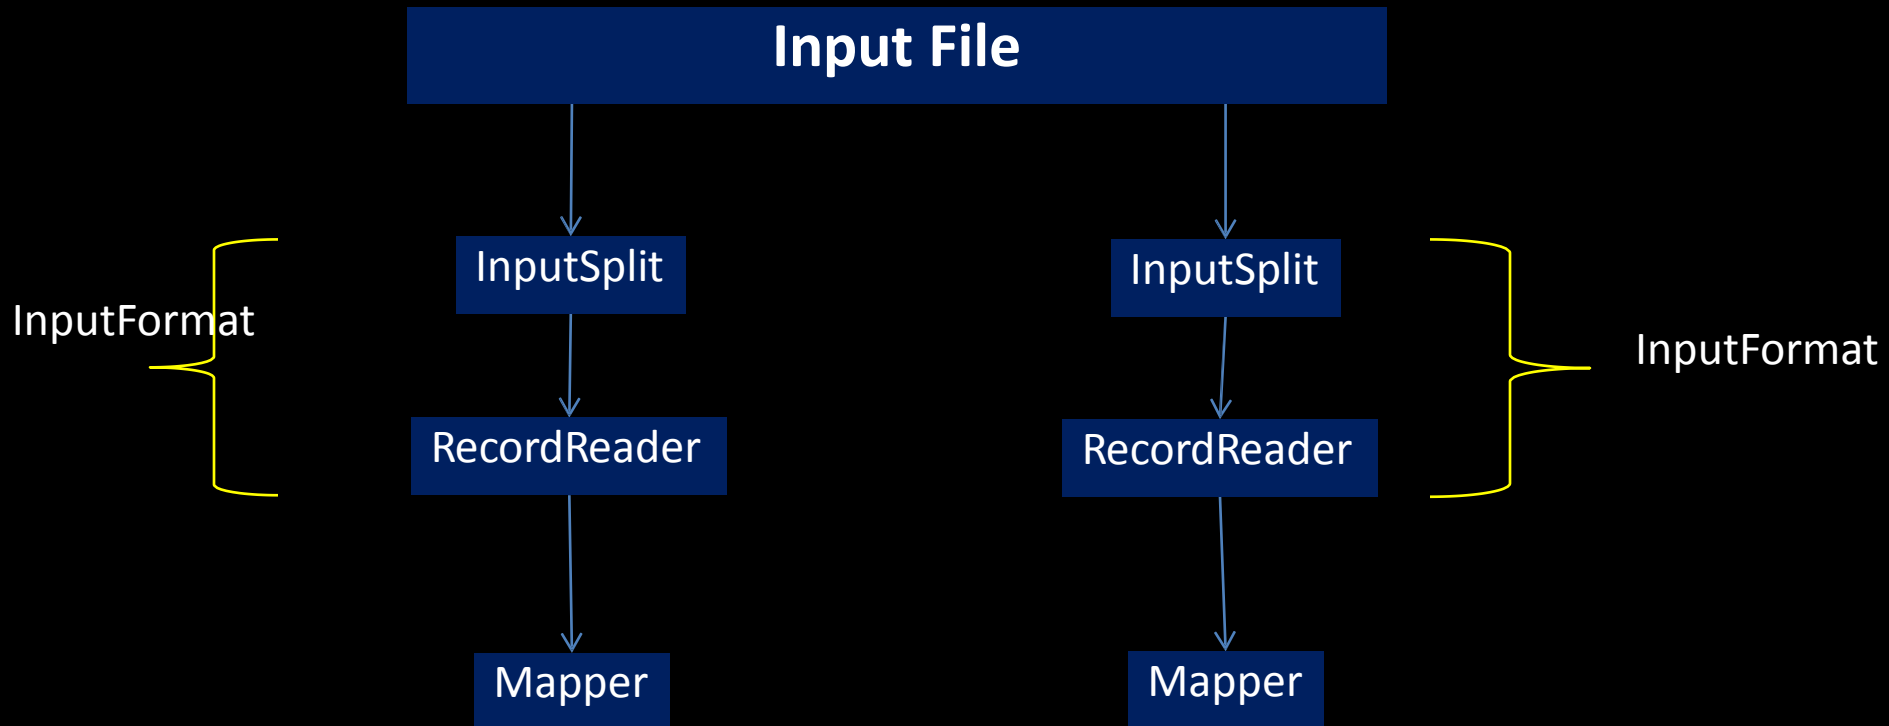
- To use other languages for writing MR
- Python, perl, ruby
- RAD, use of existing libraries
- Use stdin, stdout for input & output respectively
- TextInputFormat –default -- without key
- MR emits Key(tab)value
- No iterators as that of Java- developers have to make sure to detect change in key

```
bin/hadoop jar contrib/streaming/hadoop-streaming-1.0.4.jar -file /hadoop/hadoop-docs/mapper.py -mapper /hadoop/hadoop-docs/mapper.py -file /hadoop/hadoop-docs/reducer.py -reducer /hadoop/hadoop-docs/reducer.py -input /data_30lac.txt -output /python_out
```

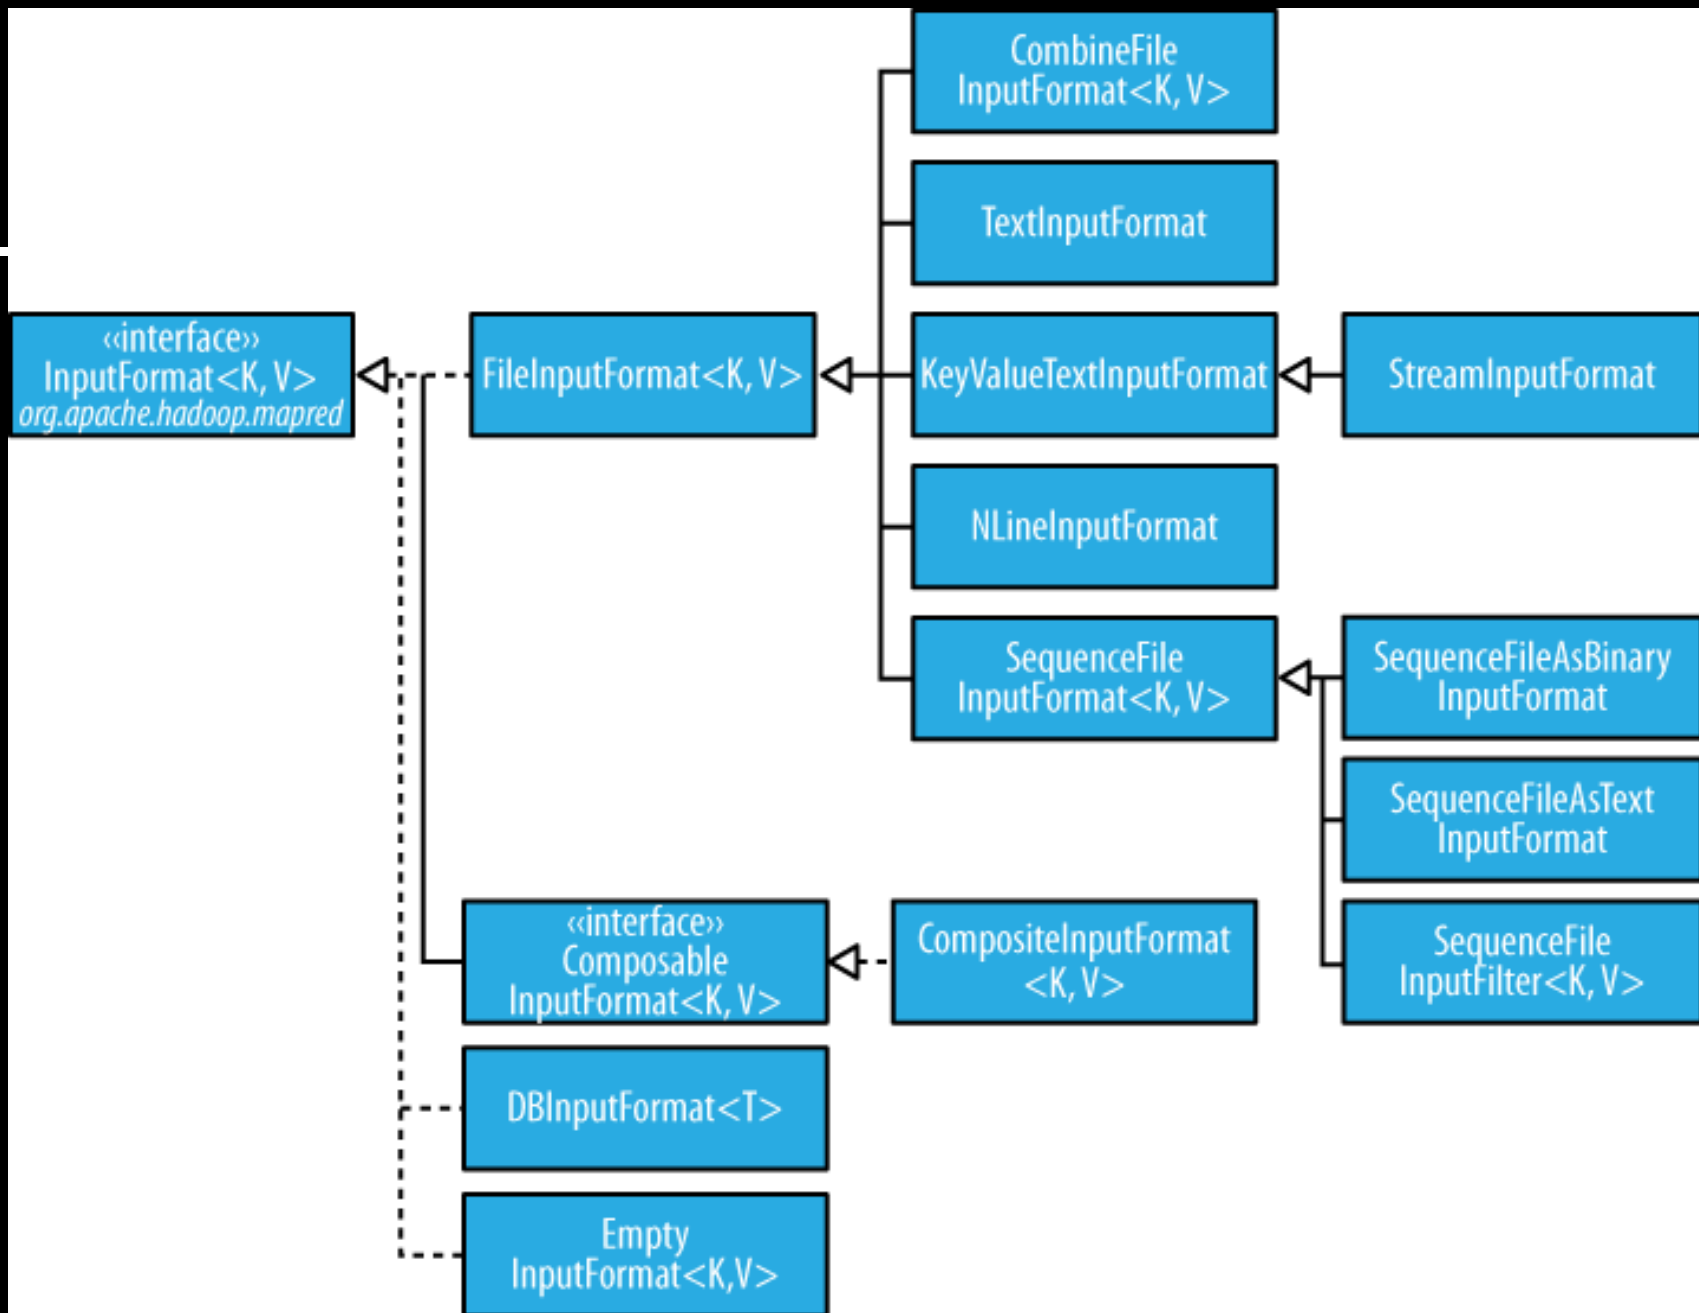
To Change Input Split Size

mapred.min.split.size	mapred.max.split.size	dfs.block.size	Split Size
1 (default)	Long.MAX_VALUE (default)	64 MB (default)	64 MB
1 (default)	Long.MAX_VALUE (default)	128 MB	128 MB
128 MB	Long.MAX_VALUE (default)	64 MB (default)	128 MB
1 (default)	10 MB	64 MB (default)	10 MB

InputFormat



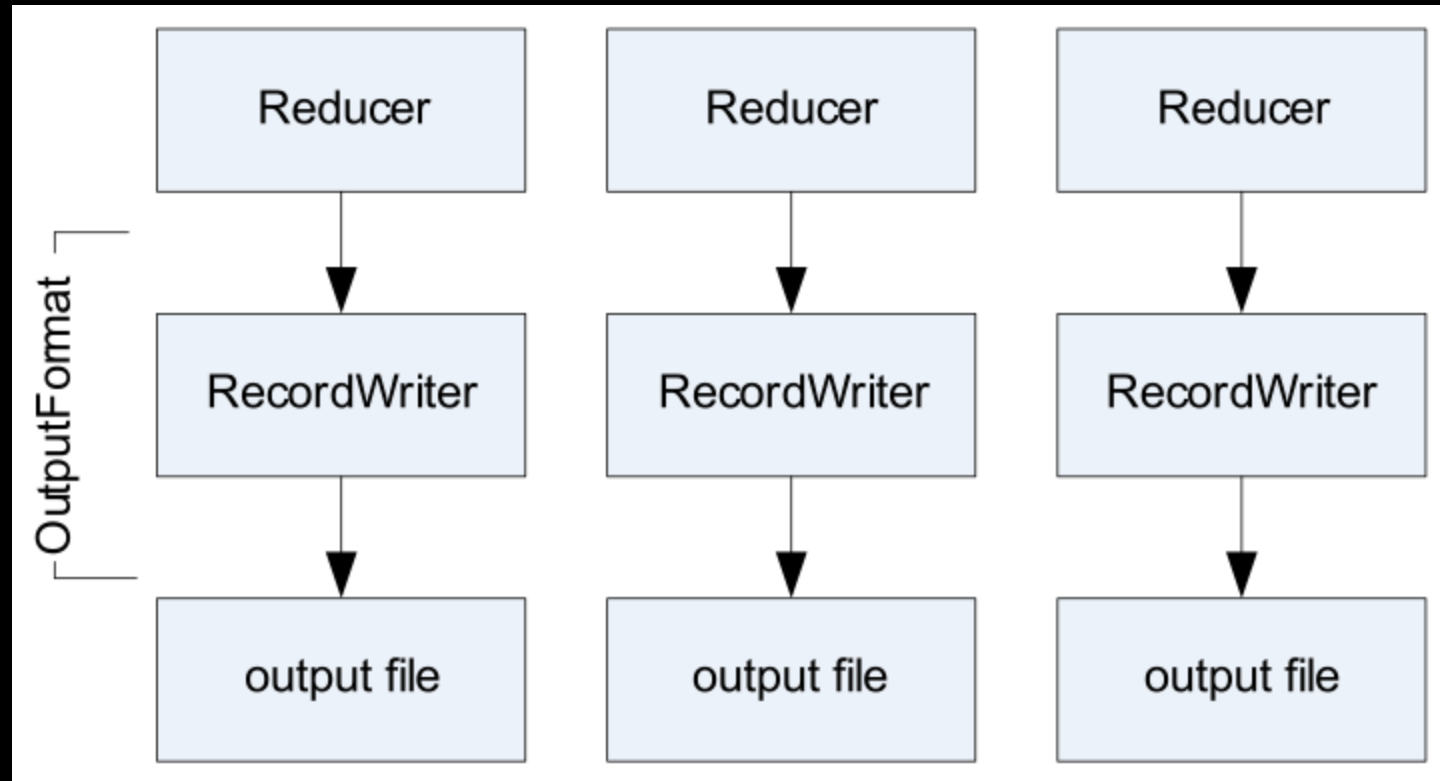
Responsible for creating inputsplit and dividing them into records

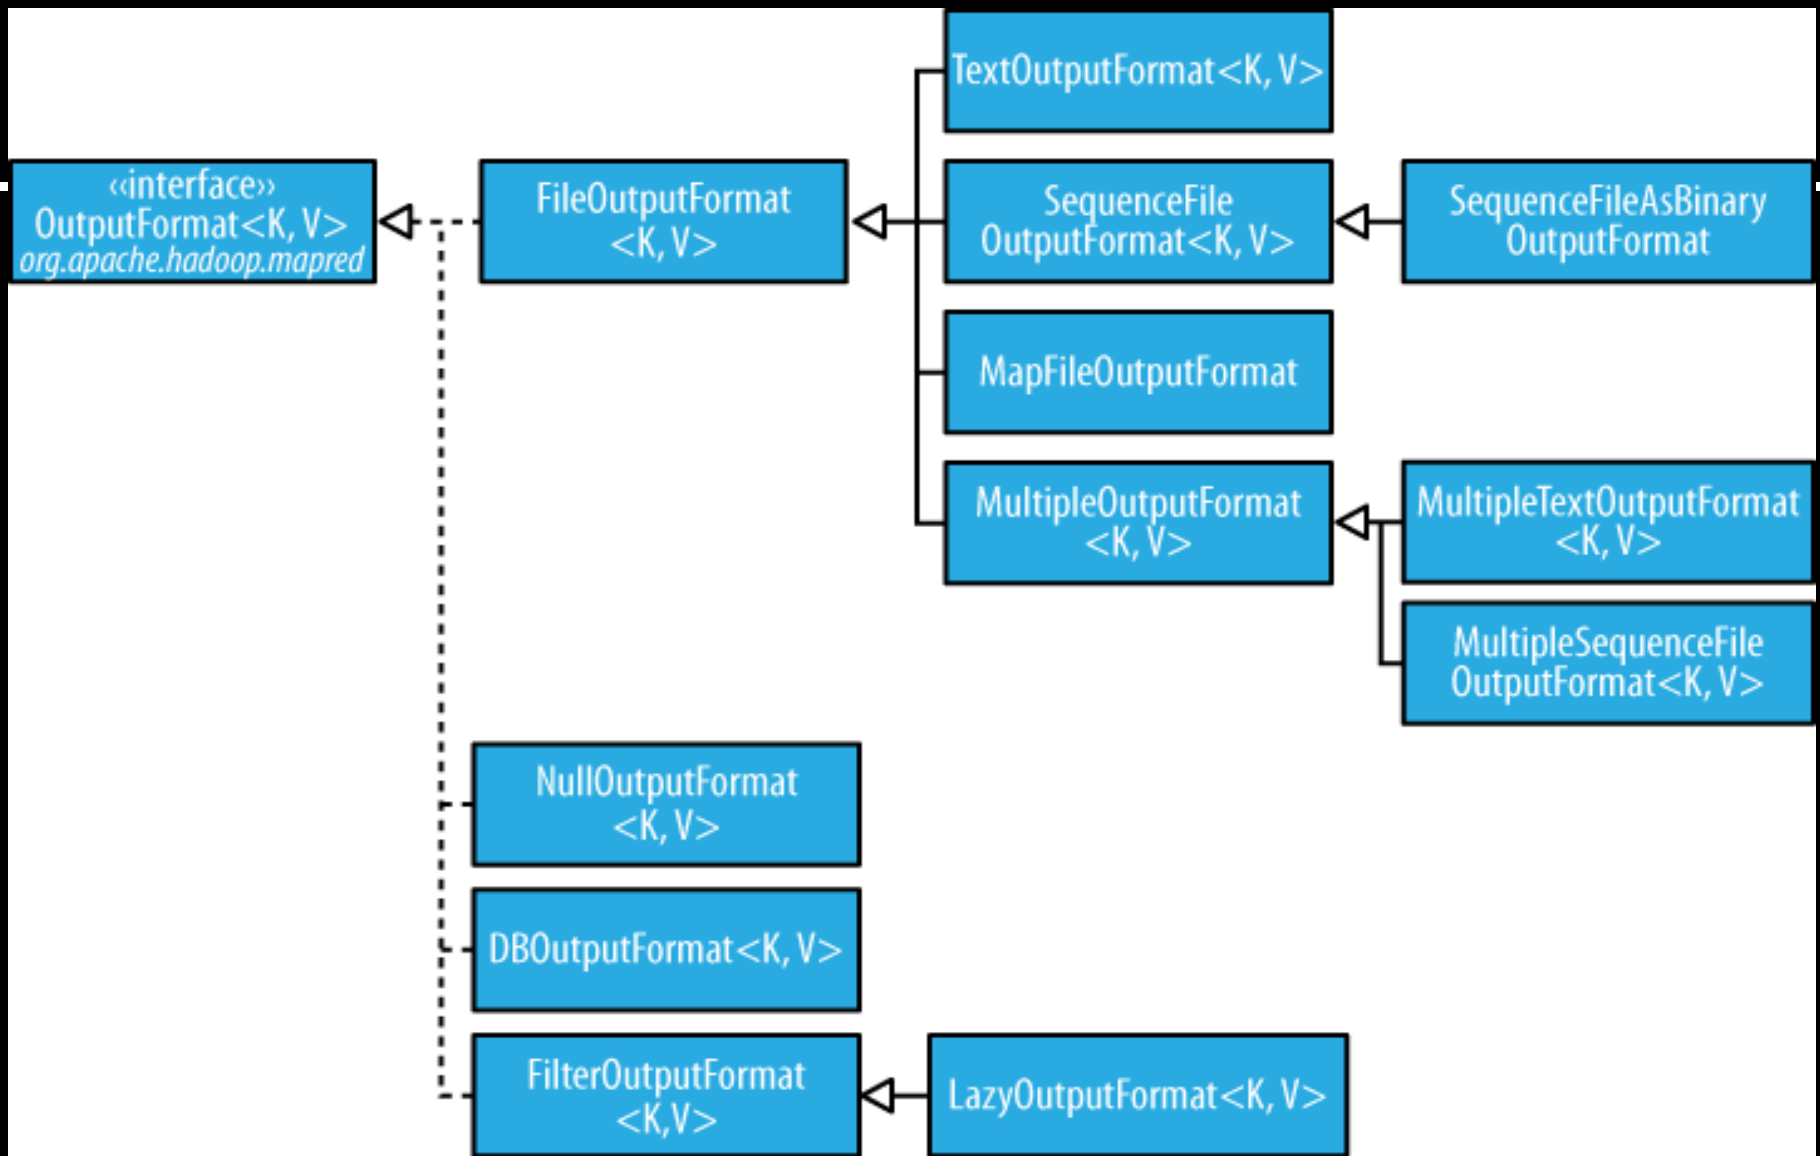


InputFormat	Key	Value
TextInputFormat	Byte offset of the line	Line contents
KeyValueInputFormat	Everything up to the first tab character	The remainder of the line
SequenceFileInputFormat	User-defined	User-defined
WholeFileInputFormat	NullWritable	file contents
NLineInputFormat	Byte offset of the line	<i>n</i> Number of lines
MultipleInputs	Per path basis	Per Path Basis
TableInputFormat (HBase)	Rowkey	Value

mapred.line.input.format.linespermap

OutPutFormat





CustomInputFormat

- Use FileInputFormat as a starting Point
 - -Extend it
- Write your own custom RecordReader
- Override getRecordReader method in FileInputFormat
- Override isSplittable if you don't want input files to be split

Some More Info

- IdentityMapper
 - mapping inputs directly to outputs
- IdentityReducer
- Performs no reduction, writing all input values directly to the output.
- Single Reducer
 - Use when complete sort order is required
- Zero Reducer
 - SetNumReduceTasks to 0
 - Output from maps will go directly to OutputFormat and disk
 - No Sorting and Shuffling

Counters

- for gathering statistics about the job
 - for quality control
 - for application level-statistics
- Classified into two counters
 - Built In Counters
 - Task counters
 - Job Counters
 - Custom Counters

Built in Counter for Job

	Counter	Map	Reduce	Total
Job Counters	SLOTS_MILLIS_MAPS	0	0	4,425,106
	Total time spent by all reduces waiting after reserving slots (ms)	0	0	0
	Total time spent by all maps waiting after reserving slots (ms)	0	0	0
	Rack-local map tasks	0	0	189
	Launched map tasks	0	0	337
	Data-local map tasks	0	0	148
	SLOTS_MILLIS_REDUCE	0	0	0
File Output Format Counters	Bytes Written	0	0	0
FileSystemCounters	HDFS_BYTES_READ	88,533	0	88,533
	FILE_BYTES_WRITTEN	33,492,676	0	33,492,676
	HDFS_BYTES_WRITTEN	13,932	0	13,932
File Input Format Counters	Bytes Read	0	0	0
	Map input records	774	0	774
	Physical memory (bytes) snapshot	34,978,086,912	0	34,978,086,912
	Spilled Records	0	0	0

CUSTOM COUNTER

```
public static enum CUSTOMCOUNTER {  
    ERROR_COUNT  
}
```

To insert the code in JOB class

To increase the counter value in mapper class

```
context.getCounter(CUSTOMCOUNTER.ERROR_COUNT).increment(1);
```

To display the output of the counter in
job after completion

```
job.getCounters().findCounter(CUSTOMCOUNTER.ERROR_COUNT).getValue()
```

Side Data Distribution

- To keep some read only data available for all the tasks
- Can be achieved using following two ways:
 - Configuration Object
 - DistributedCache

```
Configuration conf = new Configuration();  
conf.set("personName","kumar");
```

Distributed Cache

- Cache files and archive to the task nodes
- Usage
 - As
 - DistributedCache API
 - ✓ *public void addCacheFile(URL uri)*
 - ✓ *public void addCacheArchive(URL uri)*
 - ✓ *public void setCacheFiles(URL[] files)*
 - ✓ *Public void setcacheArchives(URL[] archives)*
 - ✓ *public void addFileToClassPath(Path file)*
 - ✓ *public void addArchiveToClassPath(Path archive)*
 - ✓ *public void createSymlink()*

GenericOptionsParser, Tool, & ToolRunner

- class that interprets common Hadoop command-line options
- implement the Tool interface and run your application with ToolRunner.
- GenericOptionsParser
 - *-Dproperty=value*
 - *-conf file*
 - *-fs uri*
 - *-jt host:port*
 - *-files file1 file2*
 - *-archives archive1 archive2*
 - *-libjars jar1*

Joins

- To join data from Multiple datasets
- Please try to use PIG or HIVE join if you are using text based files
- Two Varieties or approaches
 - Map-side Join
 - Reduce side Join

Map Side Join

■ Basic idea for Map-side joins:

- Load one set of data into memory, stored in an associative array
- Key of the associative array is the join key
- Map over the other set of data, and perform a lookup on the associative array using the join key
- If the join key is found, you have a successful join
- Otherwise, do nothing

Reduce Side Join

- Use the same key for the mapper output
- Can be performed in two ways
 - MultipleInputs
 - Secondary Sort

MultipleInputs.addInputPath(job, InputPath, TextInputFormat.class, CustomMapper.class);

Sorting

- In MR, keys from all mappers are sent to the reducers in sorting order
- Total Sorting can be obtained using single reducer
 - Reduce the performance
- Partial Sorting can be obtained using partition
- For benchmarking, terasort has been widely used

Secondary Sorting

- In MR sorting , the keys are sorted and not the values
- To achieve the secondary sorting
 - ✓ Make the key a composite of the natural key and the natural value
 - ✓ The sort comparator should order by the composite key, that is, the natural key *and* natural value.
 - ✓ The partitioner and grouping comparator for the composite key should
 - ✓ consider only the natural key for partitioning and grouping.

Example

Searching

Assignment

- Input
 - ✓ A set of input files containing lines
 - ✓ A pattern
- OutPut
 - ✓ Pattern with list of filename containing the pattern

Solution:

- ❖ Set the pattern using Configuration Object
- ❖ In Mapper(with TextInputFormat), verify for the pattern
- ❖ If Pattern matches, emit (pattern, filename)
- ❖ If No pattern , emit nothing
- ❖ Zero reducer

LocalJobRunner

- Only designed for simple testing of MapReduce programs
- Can't run more than one reducer
- `bin/hadoop jar job.jar com.example.wordcount -D mapred.job.tracker=local -D fs.default.name=file:/// (args)`
- *Other tips*
 - Set `keep.failed.task.files` to true
 - Use the `isolationRunner` to run just the failed tasks

THANK YOU!